

```
In [4]: import numpy as np
samples = ['The cat sat on the mat.', 'The dog ate my homework.']
token_index = {}
for sample in samples:
    for word in sample.split():
        if word not in token_index:
            token_index[word] = len(token_index) + 1
max_length = 10
results = np.zeros(shape=(len(samples),
                         max_length, max(token_index.values()) + 1))
for i, sample in enumerate(samples):
    for j, word in list(enumerate(sample.split()))[:max_length]:
        index = token_index.get(word)
        results[i, j, index] = 1
```

In [5]: results

```
Out[5]: array([[[0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.],  
                 [0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0.],  
                 [0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0.],  
                 [0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0.],  
                 [0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0.],  
                 [0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0.],  
                 [0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0.],  
                 [0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0.],  
                 [0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0.],  
                 [0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0.],  
                 [0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1.]],  
                [[0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.],  
                 [0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0.],  
                 [0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0.],  
                 [0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0.],  
                 [0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1.],  
                 [0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.],  
                 [0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.],  
                 [0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.],  
                 [0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.],  
                 [0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.]]])
```

```
In [6]: import string
samples = ['The cat sat on the mat.', 'The dog ate my homework.']
characters = string.printable
token_index = dict(zip(range(1, len(characters) + 1), characters))
max_length = 50
results = np.zeros((len(samples), max_length, max(token_index.keys()) + 1))
for i, sample in enumerate(samples):
    for j, character in enumerate(sample):
        index = token_index.get(character)
        results[i, j, index] = 1
```

In [7]: results

```
Out[7]: array([[1., 1., 1., ..., 1., 1., 1.],
   [1., 1., 1., ..., 1., 1., 1.],
   [1., 1., 1., ..., 1., 1., 1.],
   ...,
   [0., 0., 0., ..., 0., 0., 0.],
   [0., 0., 0., ..., 0., 0., 0.],
   [0., 0., 0., ..., 0., 0., 0.]],

   [[1., 1., 1., ..., 1., 1., 1.],
   [1., 1., 1., ..., 1., 1., 1.],
   [1., 1., 1., ..., 1., 1., 1.],
   ...,
   [0., 0., 0., ..., 0., 0., 0.],
   [0., 0., 0., ..., 0., 0., 0.],
   [0., 0., 0., ..., 0., 0., 0.]]])
```

In [8]:

```
from keras.preprocessing.text import Tokenizer
samples = ['The cat sat on the mat.', 'The dog ate my homework.']
tokenizer = Tokenizer(num_words=1000)
tokenizer.fit_on_texts(samples)
sequences = tokenizer.texts_to_sequences(samples)
one_hot_results = tokenizer.texts_to_matrix(samples, mode='binary')
word_index = tokenizer.word_index
print('Found %s unique tokens.' % len(word_index))
```

Using TensorFlow backend.

Found 9 unique tokens.

In [9]: one_hot_results

```
Out[9]: array([[0., 1., 1., ..., 0., 0., 0.],
   [0., 1., 0., ..., 0., 0., 0.]])
```

In [10]: one_hot_results.shape

```
Out[10]: (2, 1000)
```

In [11]:

```
from keras.layers import Embedding
embedding_layer = Embedding(1000, 64)
```

WARNING:tensorflow:From C:\ProgramData\Anaconda3\lib\site-packages\keras\backend\tensorflow_backend.py:66: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

```
In [13]: from keras.datasets import imdb
from keras import preprocessing
max_features = 10000
 maxlen = 20
(x_train, y_train), (x_test, y_test) = imdb.load_data(
num_words=max_features)
x_train = preprocessing.sequence.pad_sequences(x_train, maxlen=maxlen)
x_test = preprocessing.sequence.pad_sequences(x_test, maxlen=maxlen)
```

```
In [14]: from keras.models import Sequential
from keras.layers import Flatten, Dense
model = Sequential()
model.add(Embedding(10000, 8, input_length=maxlen))
model.add(Flatten())
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.summary()
history = model.fit(x_train, y_train,
epochs=10,
batch_size=32,
validation_split=0.2)
```

WARNING:tensorflow:From C:\ProgramData\Anaconda3\lib\site-packages\keras\backend\tensorflow_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From C:\ProgramData\Anaconda3\lib\site-packages\keras\backend\tensorflow_backend.py:4432: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.

WARNING:tensorflow:From C:\ProgramData\Anaconda3\lib\site-packages\keras\optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From C:\ProgramData\Anaconda3\lib\site-packages\keras\backend\tensorflow_backend.py:3657: The name tf.log is deprecated. Please use tf.math.log instead.

WARNING:tensorflow:From C:\ProgramData\Anaconda3\lib\site-packages\tensorflow\python\ops\nn_impl.py:180: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Model: "sequential_1"

Layer (type)	Output Shape	Param #
<hr/>		
embedding_2 (Embedding)	(None, 20, 8)	80000
<hr/>		
flatten_1 (Flatten)	(None, 160)	0
<hr/>		
dense_1 (Dense)	(None, 1)	161
<hr/>		
Total params: 80,161		
Trainable params: 80,161		
Non-trainable params: 0		

WARNING:tensorflow:From C:\ProgramData\Anaconda3\lib\site-packages\keras\backend\tensorflow_backend.py:1033: The name tf.assign_add is deprecated. Please use tf.compat.v1.assign_add instead.

Train on 20000 samples, validate on 5000 samples

Epoch 1/10

20000/20000 [=====] - 3s 142us/step - loss: 0.6679 - acc: 0.6211 - val_loss: 0.6134 - val_acc: 0.6942

Epoch 2/10

20000/20000 [=====] - 2s 114us/step - loss: 0.5353 - acc: 0.7530 - val_loss: 0.5203 - val_acc: 0.7352

Epoch 3/10

20000/20000 [=====] - 2s 109us/step - loss: 0.4576 - acc: 0.7880 - val_loss: 0.4990 - val_acc: 0.7478

Epoch 4/10

20000/20000 [=====] - 2s 109us/step - loss: 0.4205 - acc: 0.8082 - val_loss: 0.4916 - val_acc: 0.7560

Epoch 5/10

20000/20000 [=====] - 2s 109us/step - loss: 0.3953 - acc: 0.8238 - val_loss: 0.4919 - val_acc: 0.7552

Epoch 6/10

```
20000/20000 [=====] - 3s 128us/step - loss: 0.3748 -  
acc: 0.8349 - val_loss: 0.4938 - val_acc: 0.7578  
Epoch 7/10  
20000/20000 [=====] - 2s 121us/step - loss: 0.3554 -  
acc: 0.8465 - val_loss: 0.4978 - val_acc: 0.7584  
Epoch 8/10  
20000/20000 [=====] - 4s 178us/step - loss: 0.3368 -  
acc: 0.8557 - val_loss: 0.5031 - val_acc: 0.7556  
Epoch 9/10  
20000/20000 [=====] - 3s 149us/step - loss: 0.3187 -  
acc: 0.8663 - val_loss: 0.5089 - val_acc: 0.7550  
Epoch 10/10  
20000/20000 [=====] - 2s 111us/step - loss: 0.3014 -  
acc: 0.8760 - val_loss: 0.5156 - val_acc: 0.7548
```

In [15]:

```
import os  
imdb_dir='D:/Deeplearning/datasets/imdbdataset/aclImdb_v1/aclImdb'  
train_dir=os.path.join(imdb_dir,'train')  
labels=[]  
texts=[]  
for label_type in ['neg','pos']:  
    dir_name=os.path.join(train_dir,label_type)  
    for fname in os.listdir(dir_name):  
        if fname[-4:]=='.txt':  
            f=open(os.path.join(dir_name,fname),encoding='utf8')  
            texts.append(f.read())  
            f.close()  
        if label_type=='neg':  
            labels.append(0)  
        else:  
            labels.append(1)
```

```
In [16]: from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
import numpy as np
maxlen = 100
training_samples = 200
validation_samples = 10000
max_words = 10000
tokenizer = Tokenizer(num_words=max_words)
tokenizer.fit_on_texts(texts)
sequences = tokenizer.texts_to_sequences(texts)
word_index = tokenizer.word_index
print('Found %s unique tokens.' % len(word_index))
data = pad_sequences(sequences, maxlen=maxlen)
labels = np.asarray(labels)
print('Shape of data tensor:', data.shape)
print('Shape of label tensor:', labels.shape)
indices = np.arange(data.shape[0])
np.random.shuffle(indices)
data = data[indices]
labels = labels[indices]
x_train = data[:training_samples]
y_train = labels[:training_samples]
x_val = data[training_samples: training_samples + validation_samples]
y_val = labels[training_samples: training_samples + validation_samples]
```

Found 88582 unique tokens.
 Shape of data tensor: (25000, 100)
 Shape of label tensor: (25000,)

```
In [19]: glove_dir = 'D:/Deeplearning/datasets/glove.6B'
embeddings_index = {}
f = open(os.path.join(glove_dir, 'glove.6B.100d.txt'), encoding='utf')
for line in f:
    values = line.split()
    word = values[0]
    coefs = np.asarray(values[1:], dtype='float32')
    embeddings_index[word] = coefs
f.close()
print('Found %s word vectors.' % len(embeddings_index))
```

Found 400000 word vectors.

```
In [21]: embedding_dim = 100
embedding_matrix = np.zeros((max_words, embedding_dim))
for word, i in word_index.items():
    if i < max_words:
        embedding_vector = embeddings_index.get(word)
        if embedding_vector is not None:
            embedding_matrix[i] = embedding_vector
```

```
In [24]: from keras.models import Sequential
from keras.layers import Embedding, Flatten, Dense
model = Sequential()
model.add(Embedding(max_words, embedding_dim, input_length=maxlen))
model.add(Flatten())
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
<hr/>		
embedding_4 (Embedding)	(None, 100, 100)	1000000
flatten_3 (Flatten)	(None, 10000)	0
dense_4 (Dense)	(None, 32)	320032
dense_5 (Dense)	(None, 1)	33
<hr/>		
Total params: 1,320,065		
Trainable params: 1,320,065		
Non-trainable params: 0		

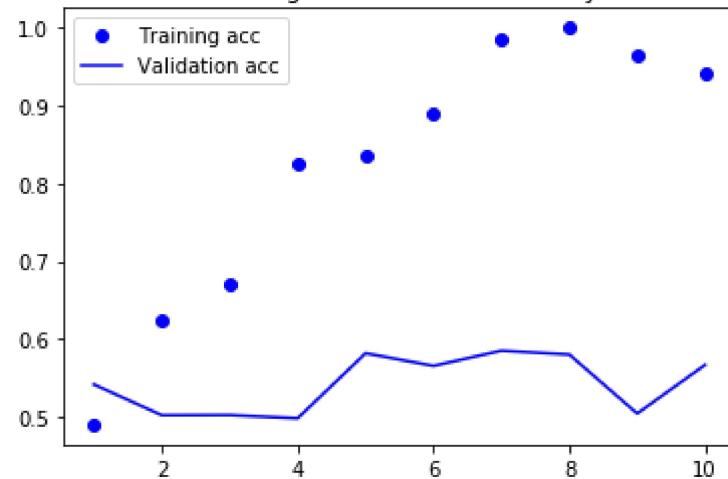
```
In [25]: model.layers[0].set_weights([embedding_matrix])
model.layers[0].trainable = False
```

```
In [26]: model.compile(optimizer='rmsprop',
loss='binary_crossentropy',
metrics=['acc'])
history = model.fit(x_train, y_train,
epochs=10,
batch_size=32,
validation_data=(x_val, y_val))
model.save_weights('pre_trained_glove_model.h5')
```

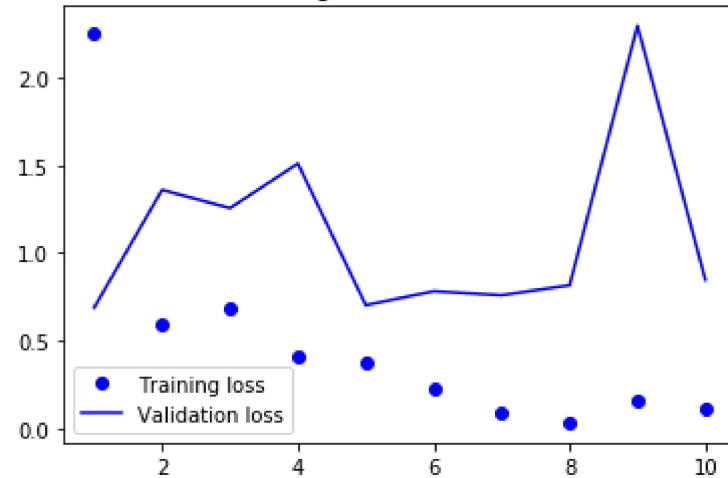
```
Train on 200 samples, validate on 10000 samples
Epoch 1/10
200/200 [=====] - 1s 7ms/step - loss: 2.2557 - acc: 0.4900 - val_loss: 0.6898 - val_acc: 0.5416
Epoch 2/10
200/200 [=====] - 1s 5ms/step - loss: 0.5913 - acc: 0.6250 - val_loss: 1.3597 - val_acc: 0.5020
Epoch 3/10
200/200 [=====] - 1s 5ms/step - loss: 0.6849 - acc: 0.6700 - val_loss: 1.2557 - val_acc: 0.5022
Epoch 4/10
200/200 [=====] - 1s 5ms/step - loss: 0.4132 - acc: 0.8250 - val_loss: 1.5097 - val_acc: 0.4981
Epoch 5/10
200/200 [=====] - 1s 6ms/step - loss: 0.3813 - acc: 0.8350 - val_loss: 0.7035 - val_acc: 0.5816
Epoch 6/10
200/200 [=====] - 1s 6ms/step - loss: 0.2271 - acc: 0.8900 - val_loss: 0.7831 - val_acc: 0.5655
Epoch 7/10
200/200 [=====] - 1s 6ms/step - loss: 0.0915 - acc: 0.9850 - val_loss: 0.7601 - val_acc: 0.5849
Epoch 8/10
200/200 [=====] - 1s 5ms/step - loss: 0.0357 - acc: 1.0000 - val_loss: 0.8167 - val_acc: 0.5800
Epoch 9/10
200/200 [=====] - 1s 6ms/step - loss: 0.1551 - acc: 0.9650 - val_loss: 2.2936 - val_acc: 0.5042
Epoch 10/10
200/200 [=====] - 1s 5ms/step - loss: 0.1095 - acc: 0.9400 - val_loss: 0.8481 - val_acc: 0.5668
```

```
In [28]: import matplotlib.pyplot as plt
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```

Training and validation accuracy



Training and validation loss



```
In [29]: from keras.models import Sequential
from keras.layers import Embedding, Flatten, Dense
model = Sequential()
model.add(Embedding(max_words, embedding_dim, input_length=maxlen))
model.add(Flatten())
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.summary()
model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['acc'])
history = model.fit(x_train, y_train,
                      epochs=10,
                      batch_size=32,
                      validation_data=(x_val, y_val))
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
<hr/>		
embedding_5 (Embedding)	(None, 100, 100)	1000000
flatten_4 (Flatten)	(None, 10000)	0
dense_6 (Dense)	(None, 32)	320032
dense_7 (Dense)	(None, 1)	33
<hr/>		
Total params: 1,320,065		
Trainable params: 1,320,065		
Non-trainable params: 0		

Train on 200 samples, validate on 10000 samples

Epoch 1/10
 200/200 [=====] - 2s 10ms/step - loss: 0.6946 - acc: 0.4550 - val_loss: 0.6920 - val_acc: 0.5219

Epoch 2/10
 200/200 [=====] - 1s 7ms/step - loss: 0.5000 - acc: 1.0000 - val_loss: 0.6927 - val_acc: 0.5269

Epoch 3/10
 200/200 [=====] - 1s 6ms/step - loss: 0.2652 - acc: 1.0000 - val_loss: 0.6932 - val_acc: 0.5248

Epoch 4/10
 200/200 [=====] - 1s 6ms/step - loss: 0.1134 - acc: 1.0000 - val_loss: 0.6939 - val_acc: 0.5326

Epoch 5/10
 200/200 [=====] - 1s 7ms/step - loss: 0.0523 - acc: 1.0000 - val_loss: 0.6979 - val_acc: 0.5302

Epoch 6/10
 200/200 [=====] - 1s 6ms/step - loss: 0.0262 - acc: 1.0000 - val_loss: 0.7007 - val_acc: 0.5326

Epoch 7/10
 200/200 [=====] - 1s 6ms/step - loss: 0.0143 - acc: 1.0000 - val_loss: 0.7039 - val_acc: 0.5363

Epoch 8/10
 200/200 [=====] - 1s 6ms/step - loss: 0.0083 - acc: 1.0000 - val_loss: 0.7089 - val_acc: 0.5344

Epoch 9/10
 200/200 [=====] - 1s 6ms/step - loss: 0.0050 - acc: 1.0000 - val_loss: 0.7128 - val_acc: 0.5334

Epoch 10/10
 200/200 [=====] - 1s 6ms/step - loss: 0.0030 - acc: 1.0000 - val_loss: 0.7166 - val_acc: 0.5362

```
In [31]: test_dir = os.path.join(imdb_dir, 'test')
labels = []
texts = []
for label_type in ['neg', 'pos']:
    dir_name = os.path.join(test_dir, label_type)
    for fname in sorted(os.listdir(dir_name)):
        if fname[-4:] == '.txt':
            f = open(os.path.join(dir_name, fname), encoding='utf')
            texts.append(f.read())
            f.close()
        if label_type == 'neg':
            labels.append(0)
        else:
            labels.append(1)
sequences = tokenizer.texts_to_sequences(texts)
x_test = pad_sequences(sequences, maxlen=maxlen)
y_test = np.asarray(labels)
```

```
In [32]: model.load_weights('pre_trained_glove_model.h5')
model.evaluate(x_test, y_test)
```

25000/25000 [=====] - 3s 103us/step

Out[32]: [0.8409046409893036, 0.57416]

```
In [33]: from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
import numpy as np
maxlen = 500
training_samples = 200
validation_samples = 10000
max_words = 10000
tokenizer = Tokenizer(num_words=max_words)
tokenizer.fit_on_texts(texts)
sequences = tokenizer.texts_to_sequences(texts)
word_index = tokenizer.word_index
print('Found %s unique tokens.' % len(word_index))
data = pad_sequences(sequences, maxlen=maxlen)
labels = np.asarray(labels)
print('Shape of data tensor:', data.shape)
print('Shape of label tensor:', labels.shape)
indices = np.arange(data.shape[0])
np.random.shuffle(indices)
data = data[indices]
labels = labels[indices]
x_train = data[:training_samples]
y_train = labels[:training_samples]
x_val = data[training_samples: training_samples + validation_samples]
y_val = labels[training_samples: training_samples + validation_samples]
```

Found 87393 unique tokens.
 Shape of data tensor: (25000, 500)
 Shape of label tensor: (25000,)

```
In [38]: embedding_dim = 100
embedding_matrix = np.zeros((max_words, embedding_dim))
for word, i in word_index.items():
    if i < max_words:
        embedding_vector = embeddings_index.get(word)
        if embedding_vector is not None:
            embedding_matrix[i] = embedding_vector
```

embedding_vector.shape

```
In [39]: embedding_vector.shape
```

```
Out[39]: (100,)
```

```
In [40]: len(word_index)
```

```
Out[40]: 87393
```

```
In [41]: from keras.models import Sequential
from keras.layers import Embedding, Flatten, Dense
model = Sequential()
model.add(Embedding(max_words, embedding_dim, input_length=maxlen))
model.add(Flatten())
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.summary()
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
<hr/>		
embedding_6 (Embedding)	(None, 500, 100)	1000000
flatten_5 (Flatten)	(None, 50000)	0
dense_8 (Dense)	(None, 32)	1600032
dense_9 (Dense)	(None, 1)	33
<hr/>		
Total params: 2,600,065		
Trainable params: 2,600,065		
Non-trainable params: 0		

```
In [42]: model.layers[0].set_weights([embedding_matrix])
model.layers[0].trainable = False
```

```
In [43]: model.compile(optimizer='rmsprop',
loss='binary_crossentropy',
metrics=['acc'])
history = model.fit(x_train, y_train,
epochs=10,
batch_size=32,
validation_data=(x_val, y_val))
model.save_weights('pre_trained_glove_model_with maxlen500.h5')
```

```
Train on 200 samples, validate on 10000 samples
Epoch 1/10
200/200 [=====] - 4s 18ms/step - loss: 5.4170 - acc: 0.5550 - val_loss: 4.5301 - val_acc: 0.4944
Epoch 2/10
200/200 [=====] - 3s 14ms/step - loss: 1.4793 - acc: 0.6100 - val_loss: 1.3307 - val_acc: 0.4963
Epoch 3/10
200/200 [=====] - 3s 14ms/step - loss: 0.5510 - acc: 0.7150 - val_loss: 0.6907 - val_acc: 0.5476
Epoch 4/10
200/200 [=====] - 3s 14ms/step - loss: 0.4283 - acc: 0.7950 - val_loss: 0.7311 - val_acc: 0.5296
Epoch 5/10
200/200 [=====] - 3s 15ms/step - loss: 0.2270 - acc: 0.9450 - val_loss: 1.1137 - val_acc: 0.5314
Epoch 6/10
200/200 [=====] - 3s 16ms/step - loss: 0.1149 - acc: 0.9850 - val_loss: 1.2881 - val_acc: 0.5105
Epoch 7/10
200/200 [=====] - 3s 14ms/step - loss: 0.7179 - acc: 0.8550 - val_loss: 0.9706 - val_acc: 0.5340
Epoch 8/10
200/200 [=====] - 3s 14ms/step - loss: 0.0422 - acc: 0.9950 - val_loss: 1.1515 - val_acc: 0.5164
Epoch 9/10
200/200 [=====] - 3s 14ms/step - loss: 0.0166 - acc: 1.0000 - val_loss: 0.8314 - val_acc: 0.5633
Epoch 10/10
200/200 [=====] - 3s 14ms/step - loss: 0.0080 - acc: 1.0000 - val_loss: 0.8109 - val_acc: 0.5699
```

```
In [45]: test_dir = os.path.join(imdb_dir, 'test')
labels = []
texts = []
for label_type in ['neg', 'pos']:
    dir_name = os.path.join(test_dir, label_type)
    for fname in sorted(os.listdir(dir_name)):
        if fname[-4:] == '.txt':
            f = open(os.path.join(dir_name, fname), encoding='utf')
            texts.append(f.read())
            f.close()
        if label_type == 'neg':
            labels.append(0)
        else:
            labels.append(1)
sequences = tokenizer.texts_to_sequences(texts)
x_test = pad_sequences(sequences, maxlen=maxlen)
y_test = np.asarray(labels)
```

```
In [46]: model.load_weights('pre_trained_glove_model_with_maxlen500.h5')
model.evaluate(x_test, y_test)
```

25000/25000 [=====] - 7s 270us/step

Out[46]: [0.8168841763877869, 0.56632]

```
In [48]: from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
import numpy as np
maxlen = 500
training_samples = 200
validation_samples = 10000
max_words = 20000
tokenizer = Tokenizer(num_words=max_words)
tokenizer.fit_on_texts(texts)
sequences = tokenizer.texts_to_sequences(texts)
word_index = tokenizer.word_index
print('Found %s unique tokens.' % len(word_index))
data = pad_sequences(sequences, maxlen=maxlen)
labels = np.asarray(labels)
print('Shape of data tensor:', data.shape)
print('Shape of label tensor:', labels.shape)
indices = np.arange(data.shape[0])
np.random.shuffle(indices)
data = data[indices]
labels = labels[indices]
x_train = data[:training_samples]
y_train = labels[:training_samples]
x_val = data[training_samples: training_samples + validation_samples]
y_val = labels[training_samples: training_samples + validation_samples]
```

Found 87393 unique tokens.
 Shape of data tensor: (25000, 500)
 Shape of label tensor: (25000,)

```
In [49]: embedding_dim = 100
embedding_matrix = np.zeros((max_words, embedding_dim))
for word, i in word_index.items():
    if i < max_words:
        embedding_vector = embeddings_index.get(word)
        if embedding_vector is not None:
            embedding_matrix[i] = embedding_vector
```

```
In [50]: from keras.models import Sequential
from keras.layers import Embedding, Flatten, Dense
model = Sequential()
model.add(Embedding(max_words, embedding_dim, input_length=maxlen))
model.add(Flatten())
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.summary()
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
<hr/>		
embedding_7 (Embedding)	(None, 500, 100)	2000000
flatten_6 (Flatten)	(None, 50000)	0
dense_10 (Dense)	(None, 32)	1600032
dense_11 (Dense)	(None, 1)	33
<hr/>		
Total params: 3,600,065		
Trainable params: 3,600,065		
Non-trainable params: 0		

```
In [51]: model.compile(optimizer='rmsprop',
loss='binary_crossentropy',
metrics=['acc'])
history = model.fit(x_train, y_train,
epochs=10,
batch_size=32,
validation_data=(x_val, y_val))
model.save_weights('pre_trained_glove_model_with_maxlen500_incft.h5')
```

```
Train on 200 samples, validate on 10000 samples
Epoch 1/10
200/200 [=====] - 4s 21ms/step - loss: 0.9553 - acc: 0.5250 - val_loss: 0.6960 - val_acc: 0.5034
Epoch 2/10
200/200 [=====] - 3s 16ms/step - loss: 0.6257 - acc: 0.5950 - val_loss: 0.7068 - val_acc: 0.5020
Epoch 3/10
200/200 [=====] - 3s 16ms/step - loss: 0.5473 - acc: 0.6100 - val_loss: 0.7165 - val_acc: 0.5026
Epoch 4/10
200/200 [=====] - 3s 17ms/step - loss: 0.5088 - acc: 0.7000 - val_loss: 0.9153 - val_acc: 0.5024
Epoch 5/10
200/200 [=====] - 4s 18ms/step - loss: 0.4739 - acc: 0.6900 - val_loss: 1.0239 - val_acc: 0.5033
Epoch 6/10
200/200 [=====] - 3s 16ms/step - loss: 0.4028 - acc: 0.7000 - val_loss: 1.0053 - val_acc: 0.5023
Epoch 7/10
200/200 [=====] - 3s 16ms/step - loss: 0.3542 - acc: 0.8000 - val_loss: 0.7608 - val_acc: 0.5029
Epoch 8/10
200/200 [=====] - 3s 16ms/step - loss: 0.3107 - acc: 0.9350 - val_loss: 0.8585 - val_acc: 0.5025
Epoch 9/10
200/200 [=====] - 3s 16ms/step - loss: 0.1588 - acc: 1.0000 - val_loss: 0.7873 - val_acc: 0.5037
Epoch 10/10
200/200 [=====] - 3s 16ms/step - loss: 0.0833 - acc: 0.9900 - val_loss: 0.8195 - val_acc: 0.4941
```

```
In [52]: from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
import numpy as np
maxlen = 500
training_samples = 10000
validation_samples = 10000
max_words = 20000
tokenizer = Tokenizer(num_words=max_words)
tokenizer.fit_on_texts(texts)
sequences = tokenizer.texts_to_sequences(texts)
word_index = tokenizer.word_index
print('Found %s unique tokens.' % len(word_index))
data = pad_sequences(sequences, maxlen=maxlen)
labels = np.asarray(labels)
print('Shape of data tensor:', data.shape)
print('Shape of label tensor:', labels.shape)
indices = np.arange(data.shape[0])
np.random.shuffle(indices)
data = data[indices]
labels = labels[indices]
x_train = data[:training_samples]
y_train = labels[:training_samples]
x_val = data[training_samples: training_samples + validation_samples]
y_val = labels[training_samples: training_samples + validation_samples]
```

Found 87393 unique tokens.
 Shape of data tensor: (25000, 500)
 Shape of label tensor: (25000,)

```
In [53]: from keras.models import Sequential
from keras.layers import Embedding, Flatten, Dense
model = Sequential()
model.add(Embedding(max_words, embedding_dim, input_length=maxlen))
model.add(Flatten())
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
<hr/>		
embedding_8 (Embedding)	(None, 500, 100)	2000000
flatten_7 (Flatten)	(None, 50000)	0
dense_12 (Dense)	(None, 32)	1600032
dense_13 (Dense)	(None, 1)	33
<hr/>		
Total params: 3,600,065		
Trainable params: 3,600,065		
Non-trainable params: 0		

```
In [54]: model.compile(optimizer='rmsprop',
loss='binary_crossentropy',
metrics=['acc'])
history = model.fit(x_train, y_train,
epochs=10,
batch_size=32,
validation_data=(x_val, y_val))
model.save_weights('pre_trained_glove_model_with_maxlen500_inctr.h5')
```

```
Train on 10000 samples, validate on 10000 samples
Epoch 1/10
10000/10000 [=====] - 37s 4ms/step - loss: 0.6997 -
acc: 0.4936 - val_loss: 0.6932 - val_acc: 0.5030
Epoch 2/10
10000/10000 [=====] - 35s 3ms/step - loss: 0.5423 -
acc: 0.7245 - val_loss: 0.8161 - val_acc: 0.5055
Epoch 3/10
10000/10000 [=====] - 35s 3ms/step - loss: 0.0966 -
acc: 0.9793 - val_loss: 1.2070 - val_acc: 0.5006
Epoch 4/10
10000/10000 [=====] - 34s 3ms/step - loss: 0.0179 -
acc: 0.9972 - val_loss: 1.5850 - val_acc: 0.5020
Epoch 5/10
10000/10000 [=====] - 34s 3ms/step - loss: 0.0185 -
acc: 0.9975 - val_loss: 1.3962 - val_acc: 0.4982
Epoch 6/10
10000/10000 [=====] - 35s 3ms/step - loss: 0.0088 -
acc: 0.9976 - val_loss: 2.0213 - val_acc: 0.4998
Epoch 7/10
10000/10000 [=====] - 35s 3ms/step - loss: 0.0087 -
acc: 0.9978 - val_loss: 1.4575 - val_acc: 0.4987
Epoch 8/10
10000/10000 [=====] - 35s 3ms/step - loss: 0.0053 -
acc: 0.9978 - val_loss: 2.0473 - val_acc: 0.4989
Epoch 9/10
10000/10000 [=====] - 35s 3ms/step - loss: 0.0063 -
acc: 0.9978 - val_loss: 2.4881 - val_acc: 0.5006
Epoch 10/10
10000/10000 [=====] - 39s 4ms/step - loss: 0.0043 -
acc: 0.9978 - val_loss: 2.4361 - val_acc: 0.4963
```

```
In [55]: from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
import numpy as np
maxlen = 100
training_samples = 12000
validation_samples = 500
max_words = 10000
tokenizer = Tokenizer(num_words=max_words)
tokenizer.fit_on_texts(texts)
sequences = tokenizer.texts_to_sequences(texts)
word_index = tokenizer.word_index
print('Found %s unique tokens.' % len(word_index))
data = pad_sequences(sequences, maxlen=maxlen)
labels = np.asarray(labels)
print('Shape of data tensor:', data.shape)
print('Shape of label tensor:', labels.shape)
indices = np.arange(data.shape[0])
np.random.shuffle(indices)
data = data[indices]
labels = labels[indices]
x_train = data[:training_samples]
y_train = labels[:training_samples]
x_val = data[training_samples: training_samples + validation_samples]
y_val = labels[training_samples: training_samples + validation_samples]
```

Found 87393 unique tokens.
 Shape of data tensor: (25000, 100)
 Shape of label tensor: (25000,)

```
In [59]: from keras.layers import Embedding, Flatten, Dense
model = Sequential()
model.add(Embedding(max_words, embedding_dim, input_length=maxlen))
model.add(Flatten())
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.summary()
```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
<hr/>		
embedding_9 (Embedding)	(None, 100, 100)	1000000
flatten_8 (Flatten)	(None, 10000)	0
dense_14 (Dense)	(None, 32)	320032
dense_15 (Dense)	(None, 1)	33
<hr/>		
Total params: 1,320,065		
Trainable params: 1,320,065		
Non-trainable params: 0		

```
In [60]: model.compile(optimizer='rmsprop',
loss='binary_crossentropy',
metrics=['acc'])
history = model.fit(x_train, y_train,
epochs=10,
batch_size=32,
validation_data=(x_val, y_val))
model.save_weights('pre_trained_glove_model_with_maxlen500_inctr12000.h5')
```

Train on 12000 samples, validate on 500 samples
Epoch 1/10
12000/12000 [=====] - 16s 1ms/step - loss: 0.6944 -
acc: 0.4932 - val_loss: 0.6930 - val_acc: 0.4840
Epoch 2/10
12000/12000 [=====] - 16s 1ms/step - loss: 0.4438 -
acc: 0.8141 - val_loss: 0.8757 - val_acc: 0.5240
Epoch 3/10
12000/12000 [=====] - 15s 1ms/step - loss: 0.0765 -
acc: 0.9819 - val_loss: 1.2733 - val_acc: 0.4960
Epoch 4/10
12000/12000 [=====] - 15s 1ms/step - loss: 0.0233 -
acc: 0.9961 - val_loss: 1.4820 - val_acc: 0.4900
Epoch 5/10
12000/12000 [=====] - 17s 1ms/step - loss: 0.0137 -
acc: 0.9966 - val_loss: 1.6801 - val_acc: 0.4980
Epoch 6/10
12000/12000 [=====] - 15s 1ms/step - loss: 0.0096 -
acc: 0.9969 - val_loss: 1.8234 - val_acc: 0.4980
Epoch 7/10
12000/12000 [=====] - 15s 1ms/step - loss: 0.0055 -
acc: 0.9970 - val_loss: 2.2648 - val_acc: 0.4940
Epoch 8/10
12000/12000 [=====] - 16s 1ms/step - loss: 0.0048 -
acc: 0.9973 - val_loss: 2.4340 - val_acc: 0.5000
Epoch 9/10
12000/12000 [=====] - 16s 1ms/step - loss: 0.0041 -
acc: 0.9972 - val_loss: 2.6214 - val_acc: 0.5100
Epoch 10/10
12000/12000 [=====] - 16s 1ms/step - loss: 0.0040 -
acc: 0.9975 - val_loss: 2.6987 - val_acc: 0.5040

```
In [63]: sequences = tokenizer.texts_to_sequences(texts)
x_test = pad_sequences(sequences, maxlen=maxlen)
y_test = np.asarray(labels)
x_test_data = x_test[:200]
y_test_data = y_test[:200]
model.load_weights('pre_trained_glove_model_with_maxlen500_inctr12000.h5')
model.evaluate(x_test_data, y_test_data)
```

200/200 [=====] - 0s 98us/step

Out[63]: [6.021943397521973, 0.46]

```
In [62]: x_test_data.shape
```

```
Out[62]: (200, 500)
```

```
In [64]: model.load_weights('pre_trained_glove_model_with_maxlen500_inctr12000.h5')
model.evaluate(x_test, y_test)
```

```
25000/25000 [=====] - 3s 105us/step
```

```
Out[64]: [5.332694325637817, 0.502]
```

```
In [ ]:
```